

Robot Localization

Javier González Jiménez

Reference Books:

- Probabilistic Robotics. S. Thrun, W. Burgard, D. Fox. MIT Press. 2001
- Simultaneous Localization and Mapping for Mobile Robots: Introduction and Methods. Juan-Antonio Fernández-Madrigal and José Luis Blanco Claraco. IGI-Global. 2013.

Content

- Least Squares Positioning
- Map registration
 - Landmark points (known correspondence)
 - Scan points: ICP (unknown correspondences)
- Filtering
 - Kalman Filter
 - Extended Kalman Filter
 - Particle Filter

Robot Localization

“Using sensory information to locate the robot in its environment is the most fundamental problem to providing a mobile robot with autonomous capabilities.” [Cox '91]

Given

Map of the environment
Sequence of sensor measurements

Wanted

Robot pose

Types of localization problems

- ***Position tracking***: the robot knows approximately where it is
- ***Global localization***: The robot has no clue where it is
- ***Kidnapped robot problem***: It thinks where it is but is wrong

Least Squares Positioning

Problem: Given a vector of measurements z that is related to the unknown pose x by the **linear observation model (o function)**

$$z = Hx$$

we want to find the “best” pose \hat{x}

Example:



Least Squares Positioning: Given $z = Hx$, find the (best) pose \hat{x}

$$z_{m \times 1} = H_{m \times n} x_{n \times 1}$$

n : pose, typically 3 or 6 unknowns

m : #observation constraints (equations)

$m=3$ $\begin{matrix} H \\ \blacksquare \\ n=3 \end{matrix}$

if $m=n$: unknowns (n) equals (independent) equations (m). For example, having 3 independent observations to compute the pose

- H (3x3) is invertible, direct solution for x : $x = H^{-1}z$

m $\begin{matrix} H \\ \blacksquare \\ n=3 \end{matrix}$

if $m < n$: more unknowns (n) than equations (m)

- The pose is **Non-observable**: infinitely many solutions for x ($H^T H$ not invertible)

m $\begin{matrix} H \\ \blacksquare \\ n=3 \end{matrix}$

if $m > n$: less unknowns (n) than equations (m)

- No exact solution since measurements z are affected by error: $z = Hx + e$
- Find x that is closest to the ideal \rightarrow **minimum square error**

$$\hat{x} = \arg \min_x \overset{\text{Cost function } \|e\|^2}{e^T e} = \arg \min_x [(z - Hx)^T (z - Hx)] = \arg \min_x \overset{\text{Cost function}}{\|z - Hx\|^2}$$

Solution: $\hat{x} = (H^T H)^{-1} H^T z$

Weighted Least Squares

- Measurements are not equally reliable: different “importance” of the noise error for each measurement $e_i = z_i - [Hx]_i$ Element i of the vector Hx

- This is modelled by a diagonal covariance matrix (no correlation between the error measurements, since they are assumed to be independent one to another)

$$Q = \begin{bmatrix} \sigma_{z_1}^2 & 0 & \dots \\ 0 & \sigma_{z_2}^2 & \dots \\ \vdots & \vdots & \ddots \end{bmatrix} \quad Q^{-1} = \begin{bmatrix} \frac{1}{\sigma_{z_1}^2} & 0 & \dots \\ 0 & \frac{1}{\sigma_{z_2}^2} & \dots \\ \vdots & \vdots & \ddots \end{bmatrix}$$

- Each element e_i of the error vector e is inversely weighted by the uncertainty in measurement z , i.e. $1/\sigma_{z_i}^2$

Cost function

$$\hat{x} = \arg \min_x e^T Q^{-1} e = \arg \min_x [(z - Hx)^T Q^{-1} (z - Hx)] = \arg \min_x \sum_{i=1}^m \frac{e_i^2}{\sigma_{z_i}^2}$$

Solution:

[Closed form]

$$\hat{x} = (H^T Q^{-1} H)^{-1} H^T Q^{-1} z$$

Best estimation

$$\Sigma_{\hat{x}} = (H^T Q^{-1} H)^{-1}$$

Uncertainty of the estimation

LS equivalent to ML Estimation if error is Gaussian

Maximum Likelihood (ML) Estimation [from lecture 2]:

Find x that maximizes the likelihood function $\mathcal{L}(x) \triangleq p(z|x)$

$$x_{ML} = \arg \max_x p(z|x)$$

If the error in the measurements is Gaussian, i.e. $e \sim N(0, Q)$

$$z = Hx + e \Rightarrow z \sim N(Hx, Q)$$

$$p(z|x) = K \cdot \exp\left\{-\frac{1}{2}[z - Hx]^T Q^{-1}[z - Hx]\right\}$$

$$\hat{x}_{ML} = \arg \max_x p(z|x) = \arg \max_x \{-[z - Hx]^T Q^{-1}[z - Hx]\}$$

$$= \arg \min_x [(z - Hx)^T Q^{-1}(z - Hx)] = \arg \min_x e^T Q^{-1}e = \hat{x}_{LS}$$

$$\hat{x}_{ML} = \hat{x}_{LS}$$

The LS and ML solutions are the same

Non-linear Least Squares

Non-linear observation (sensor) model: $z = h(x)$ (instead of $z = Hx$)

$$\hat{x} = \arg \min_x \|z - h(x)\|^2$$

These is a vector of m numbers (observations), not RVs!

No closed-form solution exists, but iterative:

Taylor expansion: $h(x) = h(x_0 + \delta) \cong h(x_0) + J_{h_0} \delta$

Jacobian of h evaluated at x_0

$$\|z - h(x)\|^2 \cong \|z - h(x_0) - J_{h_0} \delta\|^2 = \|\underbrace{(z - h(x_0))}_{\text{Error vector } e_0} - J_{h_0} \delta\|^2 = \|e_0 - J_{h_0} \delta\|^2$$

Equivalent optimization problem

$$\delta = \arg \min_{\delta} \|e_0 + J_e \delta\|^2$$

$$\delta = -(J_e^T J_e)^{-1} J_e^T e_0$$

3x1 3xn 3xm mx1

δ that makes this squared norm minimum

$$J_{h_0} = \left. \frac{dh}{dx} \right|_{x_0} = \begin{bmatrix} \frac{dh_1}{dx_1} & \dots & \frac{dh_1}{dx_n} \\ \vdots & & \vdots \\ \frac{dh_m}{dx_1} & \dots & \frac{dh_m}{dx_n} \end{bmatrix}_{x_0} = -J_e$$

mxn

Jacobian of the error vector $e = z - h(x)$ is the minus JACOBIAN of the sensor model $h(x)$

Weighted Non-linear Least Squares


Hessian and gradient of the cost function $f = \frac{1}{2}e(x)^T Q^{-1}e(x)$

$$\delta = -\underbrace{(J_e^T Q^{-1} J_e)^{-1}}_{H_f} \underbrace{J_e^T Q^{-1} e}_{\nabla f} = H_f^{-1} \nabla f \quad \left\{ \begin{array}{l} J_e^T Q^{-1} J_e = H_f \text{ Hessian matrix of } f \text{ (Fisher Information matrix)} \\ J_e^T Q^{-1} e = \nabla f \text{ Gradient vector of } f \end{array} \right.$$

Algorithm (known as Gauss-Newton):

1. Begin with an initial guess \hat{x}
2. Evaluate

$$\delta = -(J_e^T Q^{-1} J_e)^{-1} J_e^T Q^{-1} [z - h(\hat{x})]$$
3. Set $\hat{x} = \hat{x} - \delta$
4. If $\delta > \text{tolerance}$ goto 1
else stop




A note on Jacobian and Gradient

- The **gradient** is the vector formed by the partial derivatives of a **scalar (or real-value) function** ($f: \mathbb{R}^n \rightarrow \mathbb{R}$)

$$\text{GRADIENT of } f: \nabla f(x) = \begin{bmatrix} \frac{\partial f}{\partial x_1} & \dots & \frac{\partial f}{\partial x_n} \end{bmatrix} \in \mathbb{R}^n$$

Example (Wikipedia): Gradient of the function $f(x_1, x_2) = -(\cos^2 x_1 + \cos^2 x_2)^2$ depicted as a projected vector field on the bottom plane



- The **Jacobian** matrix is
 - the matrix formed by the partial derivatives of a **vector function** ($f: \mathbb{R}^n \rightarrow \mathbb{R}^m$)
 - a generalization of the **gradient** to vector functions.


$\nabla f_1(x)$: each row is a gradient

$$\text{JACOBIAN of } f: J = \begin{bmatrix} \frac{\partial f}{\partial x_1} & \dots & \frac{\partial f}{\partial x_n} \end{bmatrix} = \begin{bmatrix} \frac{df_1}{dx_1} & \dots & \frac{df_1}{dx_n} \\ \vdots & & \vdots \\ \frac{df_m}{dx_1} & \dots & \frac{df_m}{dx_n} \end{bmatrix}_{x_0} \quad J = \begin{bmatrix} \frac{\partial f_i}{\partial x_j} \end{bmatrix}$$

Least Squares Positioning

Example: Multi-lateration(also called *Range-only positioning*) in 2D


Estimate the position p of a vehicle in 2D, given distances d_i to n beacons m_i (sort of indoor GPS)



m_i : position of beacon i

$p=[x,y]^T$: robot position

The range-bearing observation model (revisited)



Observation model

Jacobian (wrt the robot pose \mathbf{x})

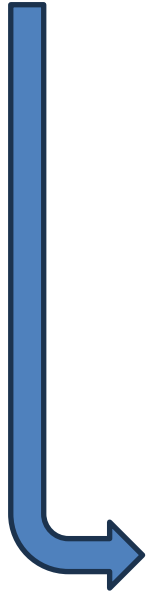
$$\mathbf{z}_i \equiv \begin{bmatrix} d_i \\ \theta_i \end{bmatrix} = h(\mathbf{x}, \mathbf{m}_i) + \mathbf{e}_i = \begin{bmatrix} \sqrt{(x_i - x)^2 + (y_i - y)^2} \\ \text{atan}\left(\frac{y_i - y}{x_i - x}\right) - \theta \end{bmatrix} + \mathbf{e}_i$$

$$\mathbf{J}_{h,\mathbf{x}} = \frac{\partial h}{\partial \{x, y, \theta\}} = \begin{bmatrix} -\frac{x_i - x}{d} & -\frac{y_i - y}{d} & 0 \\ \frac{y_i - y}{d^2} & -\frac{x_i - x}{d^2} & -1 \end{bmatrix}_{2 \times 3}$$

The observations are the true values $h(\mathbf{x}, \mathbf{m}_i)$ (from the true pose to the true landmark position) plus a sensor error \mathbf{e}_i

Example 3 beacons: Trilateration (Range-only positioning) in 2D

$$\mathbf{z} = h(\mathbf{x}) + \mathbf{e} \xrightarrow{m=3} \begin{bmatrix} h_1(x) \\ h_2(x) \\ h_3(x) \end{bmatrix} = \underbrace{\begin{bmatrix} \|m_1 - \mathbf{x}\| \\ \|m_2 - \mathbf{x}\| \\ \|m_3 - \mathbf{x}\| \end{bmatrix}}_{\substack{\text{Distance of } \mathbf{x} \text{ to} \\ \text{each landmark } m_i}} = \begin{bmatrix} [(x_1 - x)^2 + (y_1 - y)^2]^{1/2} \\ [(x_2 - x)^2 + (y_2 - y)^2]^{1/2} \\ [(x_3 - x)^2 + (y_3 - y)^2]^{1/2} \end{bmatrix}$$


 $\mathbf{e} = \mathbf{z} - h(\mathbf{x}) = \begin{bmatrix} d_1 - h_1(x) \\ d_2 - h_2(x) \\ d_3 - h_3(x) \end{bmatrix} = \begin{bmatrix} e_1 \\ e_2 \\ e_3 \end{bmatrix}$

Jacobian of $\mathbf{h} = -$ Jacobian of \mathbf{e}
 $J_h = -J_e$

Least square solution:

$$\hat{x} = \arg \min_x e^T e = \arg \min_x [e_1^2 + e_2^2 + e_3^2]$$

Example 3 beacons: Range-only positioning in 2D (cont.)

$$\begin{bmatrix} h_1(x) \\ h_2(x) \\ h_3(x) \end{bmatrix} = \begin{bmatrix} \|m_1 - \mathbf{x}\| \\ \|m_2 - \mathbf{x}\| \\ \|m_3 - \mathbf{x}\| \end{bmatrix} = \begin{bmatrix} [(x_1 - x)^2 + (y_1 - y)^2]^{1/2} \\ [(x_2 - x)^2 + (y_2 - y)^2]^{1/2} \\ [(x_3 - x)^2 + (y_3 - y)^2]^{1/2} \end{bmatrix}$$

Do not depend
on the robot
heading θ

Gradient of h_i :

$$\frac{\partial}{\partial x} [(x_i - x)^2 + (y_i - y)^2]^{1/2} = -\frac{1}{d_i} (x_i - x)$$

$$\frac{\partial}{\partial y} [(x_i - x)^2 + (y_i - y)^2]^{1/2} = -\frac{1}{d_i} (y_i - y)$$

$$\frac{\partial}{\partial \theta} [(x_i - x)^2 + (y_i - y)^2]^{1/2} = 0$$

$$J_h = \nabla h = \begin{bmatrix} -\frac{1}{d_1} (x_1 - x) & -\frac{1}{d_1} (y_1 - y) & 0 \\ -\frac{1}{d_2} (x_2 - x) & -\frac{1}{d_2} (y_2 - y) & 0 \\ -\frac{1}{d_3} (x_3 - x) & -\frac{1}{d_3} (y_3 - y) & 0 \end{bmatrix} = -J_e$$

Gradient of h_1

$J_h^T J_h$ is **not invertible** (rank = 2) \rightarrow as expected, **the heading of the robot (θ) is not observable (can not be estimated)**

Solution: Drop the unknown θ in the formulation

$$J_h = \begin{bmatrix} -\frac{1}{d_1} (x_1 - x) & -\frac{1}{d_1} (y_1 - y) \\ -\frac{1}{d_2} (x_2 - x) & -\frac{1}{d_2} (y_2 - y) \\ -\frac{1}{d_3} (x_3 - x) & -\frac{1}{d_3} (y_3 - y) \end{bmatrix} \quad \delta = (J_h^T J_h)^{-1} J_h^T [z - h(\hat{x})]$$


Example: Range-only positioning in 2D (cont.)

If just one landmark is observed, estimating $p = (x, y)$ has **infinitely many solutions!**


$$z_1 = d_1 = [(x_1 - x)^2 + (y_1 - y)^2]^{1/2} + e_d$$

The robot position will be on a circumference with radius d_1

With $e_d \sim N(0, \sigma^2)$, the robot position uncertainty will be given by a “gaussian donut” → **It is not a Gaussian distribution**



Likelihood $\mathcal{L}(x) = p(z_1|x, m_1)$ Beacon m




Bayes: $p(x|z_1, m_1) = kp(z_1|x, m_1)p(x|m_1)$

if $p(x|m_1)$ uniform $\rightarrow p(x, |z_1, m_1) = k'p(z_1|x, m_1)$

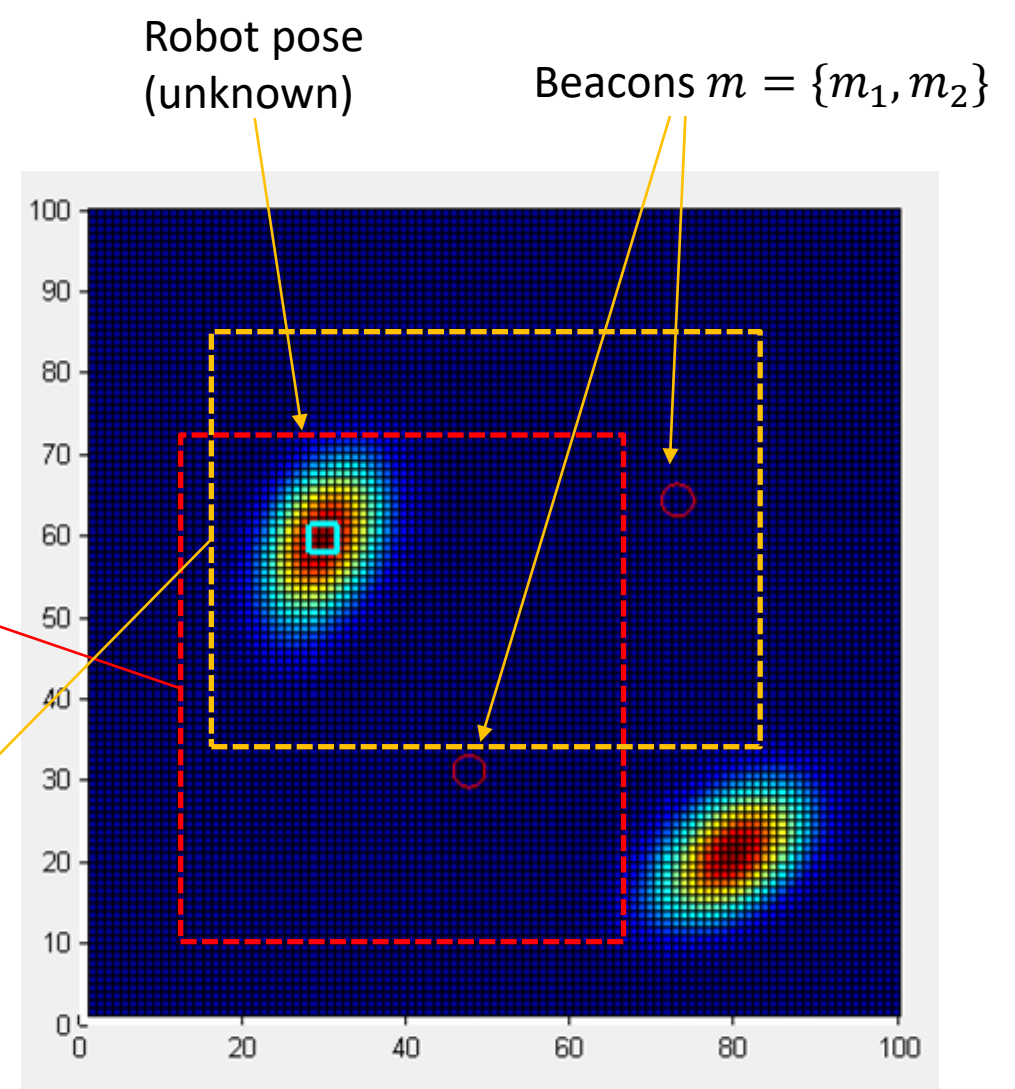

Notice: $p(z|x, m_1)$ is Gaussian BUT $\mathcal{L}(x) = p(z_1|x, m_1)$ IS NOT (“donut shape”)

Two observed range combined

First range observation
(to beacon B1)



Second range observation
(to beacon B2)



if $p(x|m)$ (prior) uniform

$$p(x|z_1, z_2, m) = k' p(z_1, z_2|x, m) =$$

Conditional independence $\rightarrow = k' p(z_1|x, m) p(z_2|x, m)$

Least Squares Positioning

Characteristics:

- Needs full observability of the state (pose) to give an estimation.
 - For example, in range-only positioning (2 unknowns: x, y), with only one observed landmark the problem becomes undetermined
 - However, using filtering techniques (i.e. Kalman filter) we can have one solution because a prior exists (explained later)
- If closed-form (non-iterative) solution exists:
 - Global minimum guaranteed → Very effective Global Localization
 - otherwise*
 - Must be solved iteratively
 - Global minimum NOT guaranteed
 - Requires a good initial guess (e.g. from the motion model)
- Recall: Least Squares is equivalent to ML Estimation if $p(z|x, m)$ is Gaussian

Map registration

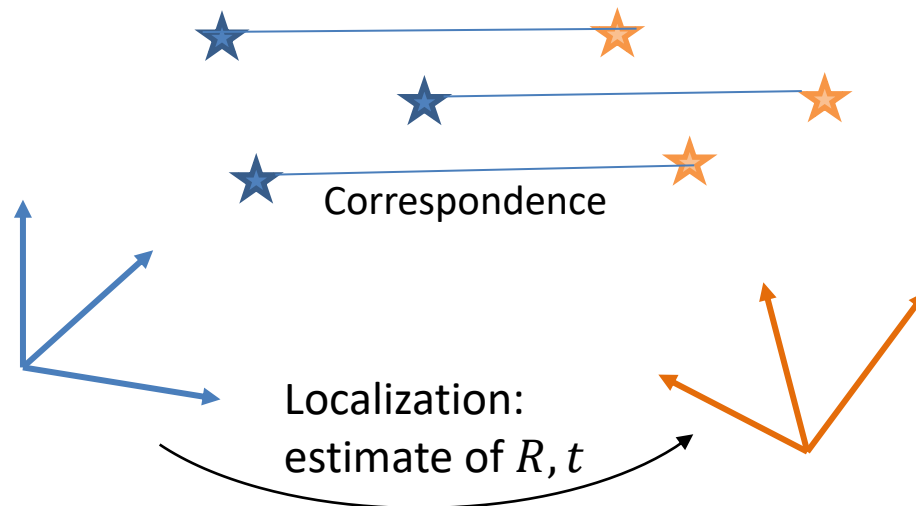
A set of observed world points are **registered (overlapped)** against ...

- A **global map** \rightarrow gives us the global pose of the observation (**global localization**)
- A **previously observed point set** \rightarrow gives us their relative pose (**odometry**)

Two types of points:


- **Landmark points:**

Correspondences are typically known
(from landmarks descriptors)




- **Scan points:**

Correspondence between points are unknown



For landmark points: known correspondences assumed

Given two sets of points $M = \{m_1, \dots, m_n\}$ and $X = \{x_1, \dots, x_n\}$ where the pair (m_i, x_i) are corresponding points, i.e. points of the same physical entity, observed from different robot poses



Error for one point:

$$E_i(R, t) = \|m_i - p \oplus x_i\|^2 = \|m_i - Rx_i - t\|^2$$

Error for n points: $E(R, t) = \sum_{i=1}^n \|m_i - Rx_i - t\|^2$

Find the **translation** t and **rotation** R that minimize the sum of the squared error $E(R, t)$

R, t can be calculated **iteratively** (e.g. Gauss-Newton) or in **closed-form** (analytically)


Scan Points Matching: Iterated Closest Point (ICP) Algorithm

Iterate until convergence:

1. For each point x_i in one set (e.g. scan points), make a pair with the **closest point** m_j in the other point set
2. Solve one step in the Gauss-Newton method for the minimization of:


$$E(R, t) = \sum_{i=1}^n \|m_j - Rx_i - t_i\|^2$$

Convergence: computed R, t in step 2 are negligible or the pairs in step 1 do not change




Closest-point matching is generally stable, but slow and requires preprocessing (KD-tree technique)

ICP: Also for 3D point clouds



[Nuechter et al., 2004]



[Jaimez & Gonzalez-Jimenez, 2015]

Discrete Kalman Filter (KF)

Kalman Filter is a method to overcome the **occasional un-observability** problem of the Least Squares approach **thanks to a model** of how the state changes

Based on the Bayes Filter (called **Markov localization**):

$$\begin{aligned} \text{Belief}(x_t) &= p(x_t | u_1, z_1 \dots, u_t, z_t) = p(x_t | u_{1:t}, z_{1:t}) \\ &= \eta \underbrace{p(z_t | x_t)}_{\text{Correction}} \underbrace{\int p(x_t | u_t, x_{t-1}) \text{Belief}(x_{t-1}) dx_{t-1}}_{\text{Prediction (prior)}} \end{aligned}$$

KF computes this equation alternating between two steps::

- **Prediction** (computes prior): $\overline{bel}(x_t) = \int p(x_t | u_t, x_{t-1}) bel(x_{t-1}) dx_{t-1}$ Total probability (sum rule)
- **Correction** (computes posterior): $bel(x_t) = \eta p(z_t | x_t) \overline{bel}(x_t)$ Bayes rule

where:

- 1) The two **process equations**, $p(x_t | u_t, x_{t-1})$ (**motion**) and $p(z_t | x_t)$ (**sensing**) are **linear**
- 2) All **pdfs** involved are **Gaussians**: $p(x_t | u_t, x_{t-1})$, $p(z_t | x_t)$, $bel(x_0)$

RECALL 1 Bayes filter:

$$\text{Belief}(x_t) = p(x_t | u_1, z_1, \dots, u_t, z_t) = p(x_t | u_{1:t}, z_{1:t})$$

$$z_{1:t} = z_1, \dots, z_t$$

$$u_{1:t} = u_1, \dots, u_t$$

Bayes $= \eta \ p(z_t | x_t, \cancel{u_{1:t}}, \cancel{z_{1:t-1}}) \ p(x_t | u_{1:t}, z_{1:t-1}) \leftarrow \text{without } z_t$

Markov $= \eta \ p(z_t | x_t) \ p(x_t | u_{1:t}, z_{1:t-1}) \quad p(x_t) = \int p(x_t, x_{t-1}) dx_{t-1} = \int p(x_t | x_{t-1}) p(x_{t-1}) dx_{t-1}$

Total prob. $= \eta \ p(z_t | x_t) \int p(x_t | \cancel{u_{1:t}}, \cancel{z_{1:t-1}}, x_{t-1}) p(x_{t-1} | u_{1:t}, z_{1:t-1}) dx_{t-1}$


Markov $= \eta \ p(z_t | x_t) \int p(x_t | u_t, x_{t-1}) p(x_{t-1} | u_{1:t}, z_{1:t-1}) dx_{t-1}$

Markov $= \eta \ p(z_t | x_t) \int p(x_t | u_t, x_{t-1}) \boxed{p(x_{t-1} | u_{1:t-1}, z_{1:t-1})} dx_{t-1}$ We don't apply Markov here because it is not needed. This is the definition of $\text{Belief}(x_{t-1})$

$$= \eta \ p(z_t | x_t) \int p(x_t | u_t, x_{t-1}) \text{Belief}(x_{t-1}) dx_{t-1}$$

RECALL 2: How Gaussians look like

$$p(\mathbf{x}) = N(\mathbf{x}; \boldsymbol{\mu}, \boldsymbol{\Sigma}) = \frac{1}{(2\pi)^{d/2} |\boldsymbol{\Sigma}|^{1/2}} e^{-\frac{1}{2}(\mathbf{x}-\boldsymbol{\mu})^t \boldsymbol{\Sigma}^{-1}(\mathbf{x}-\boldsymbol{\mu})}$$



RECALL 3:

- Product
- Marginalization/conditioning
- Linear transformation

} of Gaussians pdf's is Gaussian



We stay in the “Gaussian world” as long as we start with Gaussians and perform only linear transformations.

$$\left. \begin{array}{l} X \sim N(\mu, \Sigma) \\ Y = AX + B \end{array} \right\} \Rightarrow Y \sim N(A\mu + B, A\Sigma A^T)$$

$$\left. \begin{array}{l} X_1 \sim N(\mu_1, \Sigma_1) \\ X_2 \sim N(\mu_2, \Sigma_2) \end{array} \right\} \Rightarrow p(X_1) \cdot p(X_2) \sim N(\Sigma_{12}(\Sigma_1^{-1}\mu_1 + \Sigma_2^{-1}\mu_2), \quad \Sigma_{12} = (\Sigma_1^{-1} + \Sigma_2^{-1})^{-1})$$

KF: Let's go into the details

Kalman Filter is a general tool to estimate the state x of a process that is governed by the **linear** stochastic **model**:

$$x_t = A_t x_{t-1} + B_t u_t + \varepsilon_t$$

$$\varepsilon_t \sim N(0, R_t)$$

Prediction phase

covariance of the linear motion model $B_t u_t$

with a **linear** measurement equation

$$z_t = C_t x_t + \delta_t$$

$$\delta_t \sim N(0, Q_t)$$

Correction phase

covariance of the linear measurement model $C_t x_t$

A_t Matrix (nxn) that describes how the state x evolves from $t-1$ to t without motion command (u_t)

B_t Matrix (nx1) that describes how the motion command u_t (1x1) changes the state x from $t-1$ to t

C_t Matrix (kxn) that describes how the state x_t is related to an observation z_t .


$\varepsilon_t \delta_t$ Random variables representing **the motion and measurement noise**. They are assumed to be independent and normally distributed with covariance R_t and Q_t , respectively.

Kalman Filter from Bayes filter

$$\begin{array}{c}
 \text{Posterior} \qquad \text{Likelihood} \qquad \overline{Bel}(x_t) \text{ or prior} \\
 \text{Blue oval} \quad \text{Pink oval} \quad \text{Yellow oval} \\
 Bel(x_t) = \eta \int p(z_t | x_t) p(x_t | u_t, x_{t-1}) Bel(x_{t-1}) dx_{t-1} \\
 \downarrow \qquad \downarrow \qquad \downarrow \qquad \downarrow \qquad \downarrow \\
 N(x_t; \mu_t, \Sigma_t) \quad N(z_t; C_t x_t, Q_t) \quad N(x_t; \bar{\mu}_t, \bar{\Sigma}_t) \quad \sim N(x_{t-1}; \mu_{t-1}, \Sigma_{t-1}) \\
 \qquad \qquad \qquad \sim N(x_t; A_t x_{t-1} + B_t u_t, R_t)
 \end{array}$$

Motion model: Distribution over poses when executing the noisy motion command u_t and its pose is x_{t-1}

Distribution over poses x_{t-1}



Posterior $Bel(x_t)$ *Likelihood* $p(z_t | x_t)$ $\overline{Bel}(x_t)$ or prior

$$Bel(x_t) = \eta p(z_t | x_t) \int p(x_t | u_t, x_{t-1}) Bel(x_{t-1}) dx_{t-1}$$

$N(x_t; A_t x_{t-1} + B_t u_t, R_t)$ $N(x_{t-1}; \mu_{t-1}, \Sigma_{t-1})$

Motion model: Distribution over poses when executing the noisy motion command u_t and its pose is x_{t-1}

Distribution over poses x_{t-1}

Prediction phase (Prior computation):

$$\overline{Bel}(x_t) = \eta \int \exp \left\{ -\frac{1}{2} (x_t - A_t x_{t-1} - B_t u_t)^T R_t^{-1} (x_t - A_t x_{t-1} - B_t u_t) \right\} \exp \left\{ -\frac{1}{2} (x_{t-1} - \mu_{t-1})^T \Sigma_{t-1}^{-1} (x_{t-1} - \mu_{t-1}) \right\} dx_{t-1}$$

This is another Gaussian, since marginalization and product of Gaussians is Gaussian:

$$\overline{Bel}(x_t) = p(x_t) = \int p(x_t, x_{t-1}) dx_{t-1} = \int p(x_t | x_{t-1}) p(x_{t-1}) dx_{t-1}$$

u_t omitted for clarity

$$\overline{Bel}(x_t) = N(x_t; \bar{\mu}_t, \bar{\Sigma}_t) \begin{cases} \bar{\mu}_t = A_t \mu_{t-1} + B_t u_t \\ \bar{\Sigma}_t = A_t \Sigma_{t-1} A_t^T + B_t \Sigma_{u_t} B_t^T R_t \end{cases}$$

$$\begin{array}{ccc}
 \text{Posterior} & \text{Likelihood} & \overline{Bel}(x_t) \text{ or prior} \\
 \text{Bel}(x_t) = \eta \cdot p(z_t | x_t) \int p(x_t | u_t, x_{t-1}) \overline{Bel}(x_{t-1}) dx_{t-1} & & \\
 \downarrow & \downarrow & \downarrow \\
 N(x_t; \mu_t, \Sigma_t) & N(z_t; C_t x_t, Q_t) & N(x_t; \bar{\mu}_t, \bar{\Sigma}_t)
 \end{array}$$

Correction phase (posterior computation):

$$z_t = C_t x_t + \delta_t \quad \Rightarrow \quad p(z_t | x_t) = N(z_t; C_t x_t, Q_t) \quad \text{Observation model}$$

$$Bel(x_t) = \eta \exp \left\{ -\frac{1}{2} (z_t - C_t x_t)^T Q_t^{-1} (z_t - C_t x_t) \right\} \exp \left\{ -\frac{1}{2} (x_t - \bar{\mu}_t)^T \bar{\Sigma}_t^{-1} (x_t - \bar{\mu}_t) \right\}$$

This is the multiplication of two Gaussians, which is also Gaussian

$$Bel(x_t) = \begin{cases} \mu_t = \bar{\mu}_t + K_t (z_t - C_t \bar{\mu}_t) \\ \Sigma_t = (I - K_t C_t) \bar{\Sigma}_t \end{cases} \quad \text{with gain} \quad K_t = \bar{\Sigma}_t C_t^T (C_t \bar{\Sigma}_t C_t^T + Q_t)^{-1}$$

Algorithm **Kalman_filter** ($\mu_{t-1}, \Sigma_{t-1}, u_t, z_t$):

Prediction:

$$\bar{\mu}_t = A_t \mu_{t-1} + B_t u_t$$

$$\bar{\Sigma}_t = A_t \Sigma_{t-1} A_t^T + R_t$$

More uncertainty: The predicted covariance is the sum of those of the two RVs involved: R_t for $B_t u_t$ and $A_t \Sigma_{t-1} A_t^T$ for x_{t-1}

Correction:

$$K_t = \bar{\Sigma}_t C_t^T (C_t \bar{\Sigma}_t C_t^T + Q_t)^{-1}$$

Covariance of the **Innovation**

$$\mu_t = \bar{\mu}_t + K_t (z_t - C_t \bar{\mu}_t)$$

This is called: **Innovation**: error between the observation and the prediction of the landmark

$$\Sigma_t = (I - K_t C_t) \bar{\Sigma}_t$$

Less uncertainty: $\det(\Sigma_t) < \det(\bar{\Sigma}_t)$

Return μ_t, Σ_t

Discrete Kalman Filter

Summary

- **Highly efficient**: Polynomial complexity in dimensionality of the measurement k and state dimensionality n ($n=3$ for 2D pose):
 $O(k^{2.376} + n^2)$
- Gives an **optimal estimate** for linear Gaussian systems
- Unfortunately, most robotics models (motion and sensing) are **nonlinear** 😞, and hence, KF can not be applied ...

BUT we can **linearize the motion and sensing model** → Extended KF

Extended Kalman Filter

- EKF is an adaptation of KF to **Nonlinear Dynamic Systems**

Model transition (**prediction**): $\bar{x}_t = x_{t-1} \oplus u_t + \varepsilon_t$ (Instead of $\bar{x}_t = A_t x_{t-1} + B_t u_t + \varepsilon_t$)

Observation function (**correction**): $z_t = h(x_t) + \delta_t$ (Instead of $z_t = C_t x_t + \delta_t$)



PROBLEM: Transformed Gaussians are **no longer Gaussians**




The robot pose estimate can not be represented in a parametric form (mean, covariance)

Watch out: \bar{x}_t is NOT the mean of x_t , but the **predicted** x_t (before the observation z_t)

Linearity Assumption Revisited (lecture 2)








$$\left. \begin{array}{l} X \sim N(\mu, \Sigma) \\ Y = AX + B \end{array} \right\} \Rightarrow Y \sim N(A\mu + B, A\Sigma A^T)$$




Linearity Assumption Revisited (lecture 2)

The pdf of $y=f(x)$ is not well approximated by a Gaussian



Large variance of x



Small variance of x

For the same linearization of the function $f(x)$, **much more error in the gaussian approximation when covariance large!!**

EKF Linearization: First Order Taylor Expansion

- Prediction:** $\bar{x}_t = x_{t-1} \oplus u_t + \varepsilon_t = g(x_{t-1}, u_t) + \varepsilon_t$

$$g(x_{t-1}, u_t) \approx g(\mu_{t-1}, u_t) + \underbrace{\frac{\partial g(\mu_{t-1}, u_t)}{\partial x_{t-1}}}_{\text{Jacobian } G} \underbrace{(x_{t-1} - \mu_{t-1})}_{\Delta x_{t-1}} = g(\mu_{t-1}, u_t) + G \Delta x_{t-1}$$

Annotations for the first equation:

- mean of x_{t-1} (points to μ_{t-1})
- The Jacobian is evaluated here, the mean of x_{t-1} (points to the Jacobian term)
- Linearization for x_{t-1} at μ_{t-1} (points to the entire approximation)

$$\bar{x}_t = g(x_{t-1}, u_t) + \varepsilon_t \approx \underbrace{g(\mu_{t-1}, u_t)}_{\mu_{t-1} \oplus u_t} + G \Delta x_{t-1} + \varepsilon_t$$


Annotations for the second equation:

- Random variables $\Delta x_{t-1} \sim N(0, \Sigma_{t-1})$ and $\varepsilon_t \sim N(0, R_t)$ (points to the noise terms)


Parameters of the Normal distribution of \bar{x}_t :

$$\bar{x}_t \sim N(\bar{\mu}_t, \bar{\Sigma}_t) = N(\mu_{t-1} \oplus u_t, G_t \Sigma_{t-1} G_t^T + R_t)$$


EKF Prediction Step



Motion of a **known pose** μ_{t-1} by $u_t \sim N(\bar{u}_t, R_t)$



Motion of the pose $x_{t-1} \sim N(\mu_{t-1}, \Sigma_{t-1})$ by a **known motion** u_t



Motion of the pose $x_{t-1} \sim N(\mu_{t-1}, \Sigma_{t-1})$ by $u_t \sim N(\bar{u}_t, R_t)$

EKF Linearization: First Order Taylor Expansion

- Correction:**

$$h(x_t) \approx h(\bar{\mu}_t) + \frac{\partial h(\bar{\mu}_t)}{\partial x_t} (x_t - \bar{\mu}_t) = h(\bar{\mu}_t) + H_t (x_t - \bar{\mu}_t)$$

$H_t = \frac{\partial h(\bar{\mu}_t)}{\partial x_t}$

Predicted mean (computed in the prediction step: $\mu_{t-1} \oplus u_t$)

Likelihood

Prior (Prediction)

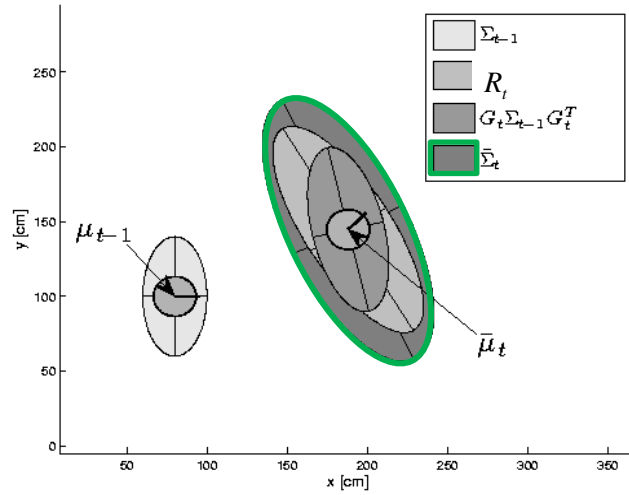
$$Bel(x_t) = \eta \exp\left\{-\frac{1}{2}(z_t - h(x_t))^T Q_t^{-1}(z_t - h(x_t))\right\} \exp\left\{-\frac{1}{2}(x_t - \bar{\mu}_t)^T \bar{\Sigma}_t^{-1}(x_t - \bar{\mu}_t)\right\}$$

Multiplication of two Gaussians = another Gaussian with a smaller Covariance


$$Bel(x_t) = N(x_t; \mu_t, \Sigma_t) \quad \left\{ \begin{array}{l} \mu_t = \bar{\mu}_t + K_t(z_t - h(\bar{\mu}_t)) \\ \Sigma_t = (I - K_t H_t) \bar{\Sigma}_t \\ K_t = \bar{\Sigma}_t H_t^T (H_t \bar{\Sigma}_t H_t^T + Q_t)^{-1} \end{array} \right.$$

EKF Correction Step

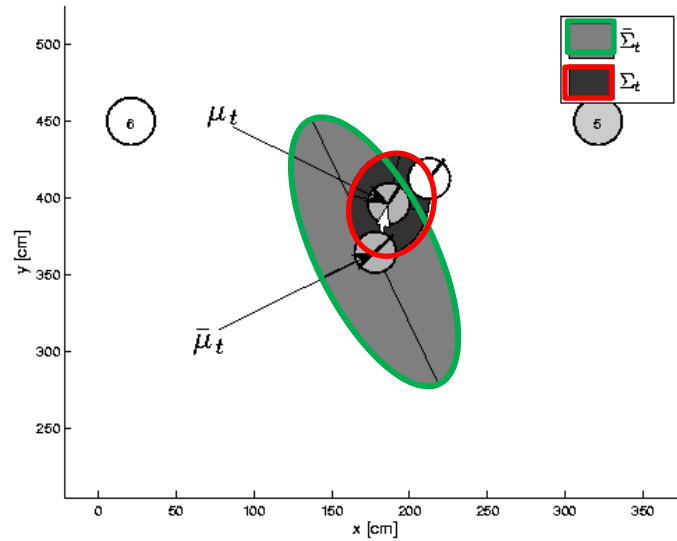
Prediction estimate



Take observation z_t



Correction estimate



Correction

Extended Kalman Filter

Extended_Kalman_filter ($\mu_{t-1}, \Sigma_{t-1}, u_t, z_t$):

Prediction:

1. $\bar{\mu}_t = g(\mu_{t-1}, u_t) = \mu_{t-1} \oplus u_t$
2. $\bar{\Sigma}_t = G_t \Sigma_{t-1} G_t^T + R_t$

Correction:

1. $K_t = \bar{\Sigma}_t H_t^T (H_t \bar{\Sigma}_t H_t^T + Q_t)^{-1}$
2. $\mu_t = \bar{\mu}_t + K_t (z_t - h(\bar{\mu}_t))$
3. $\Sigma_t = (I - K_t H_t) \bar{\Sigma}_t$

Return μ_t, Σ_t

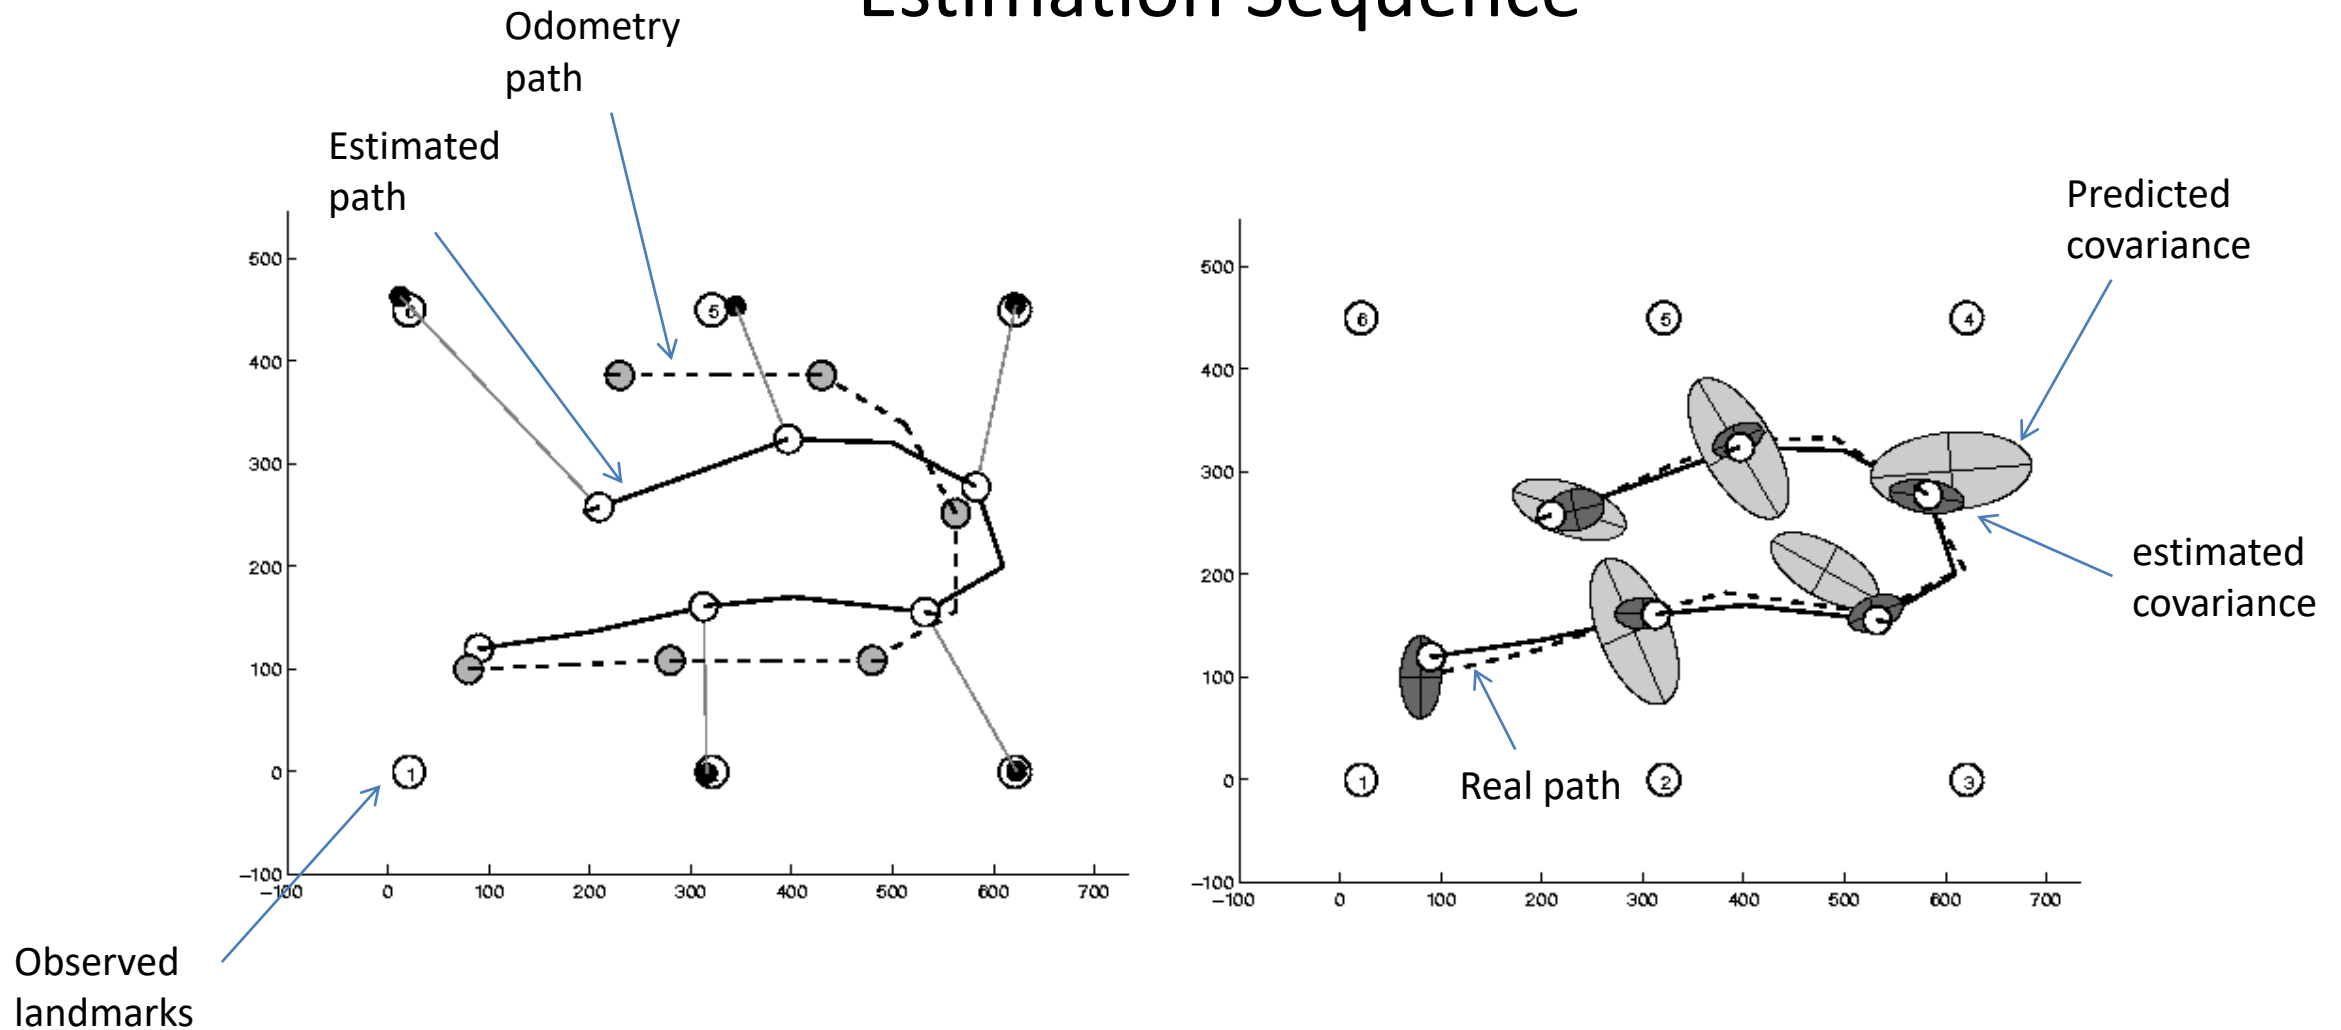
Jacobians

$$G_t = \frac{\partial g(\mu_{t-1}, u_t)}{\partial x_{t-1}}$$

$$H_t = \frac{\partial h(\bar{\mu}_t)}{\partial x_t}$$

Extended Kalman Filter

Estimation Sequence




We'll program this example in the practical sessions

Particle Filter


- Technique for implementing recursive Bayesian filter by **Monte Carlo sampling**
- **Idea:** to represent the posterior density by a set of k **random samples (particles)**. The robot pose follows the distribution given by the samples

Example for a Gaussian pdf:

1D



2D



Particle Filter

- Still an implementation of the Bayes Filter

$$Bel(x_t) = \eta p(z_t | x_t) \overline{Bel}(x_t) = \eta p(z_t | x_t) \overline{Bel}(g(u_t, x_{t-1}))$$

$$x_t = g(u_t, x_{t-1})$$

Particles x_t^i drawn from *prior*

Importance factor (weight) for x_t^i : $w_t^i \propto p(z_t | x_t)$

- Bayes Filter for pdf given by samples

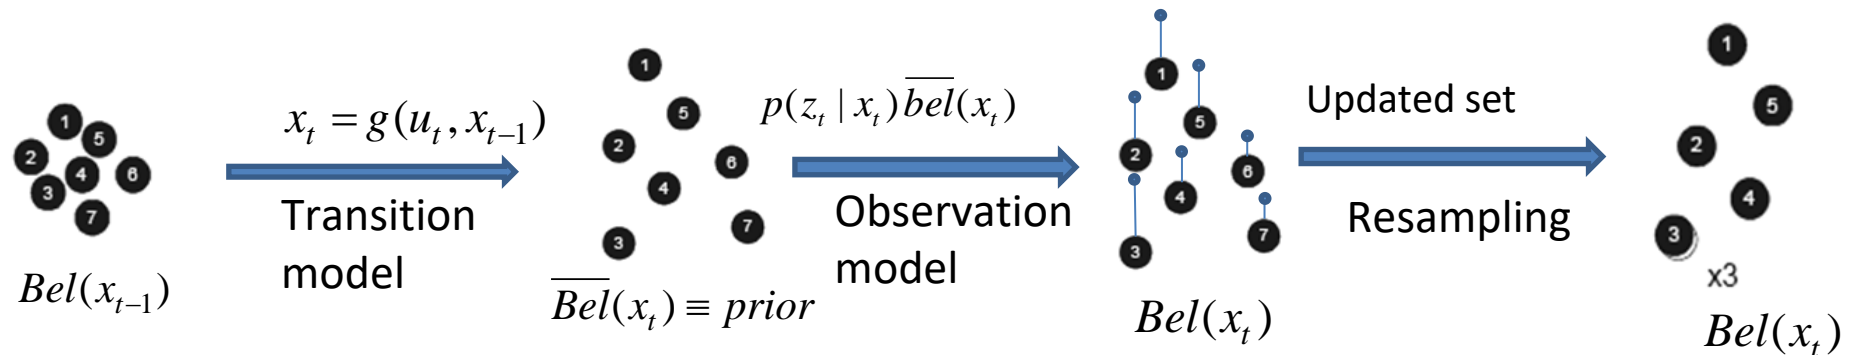
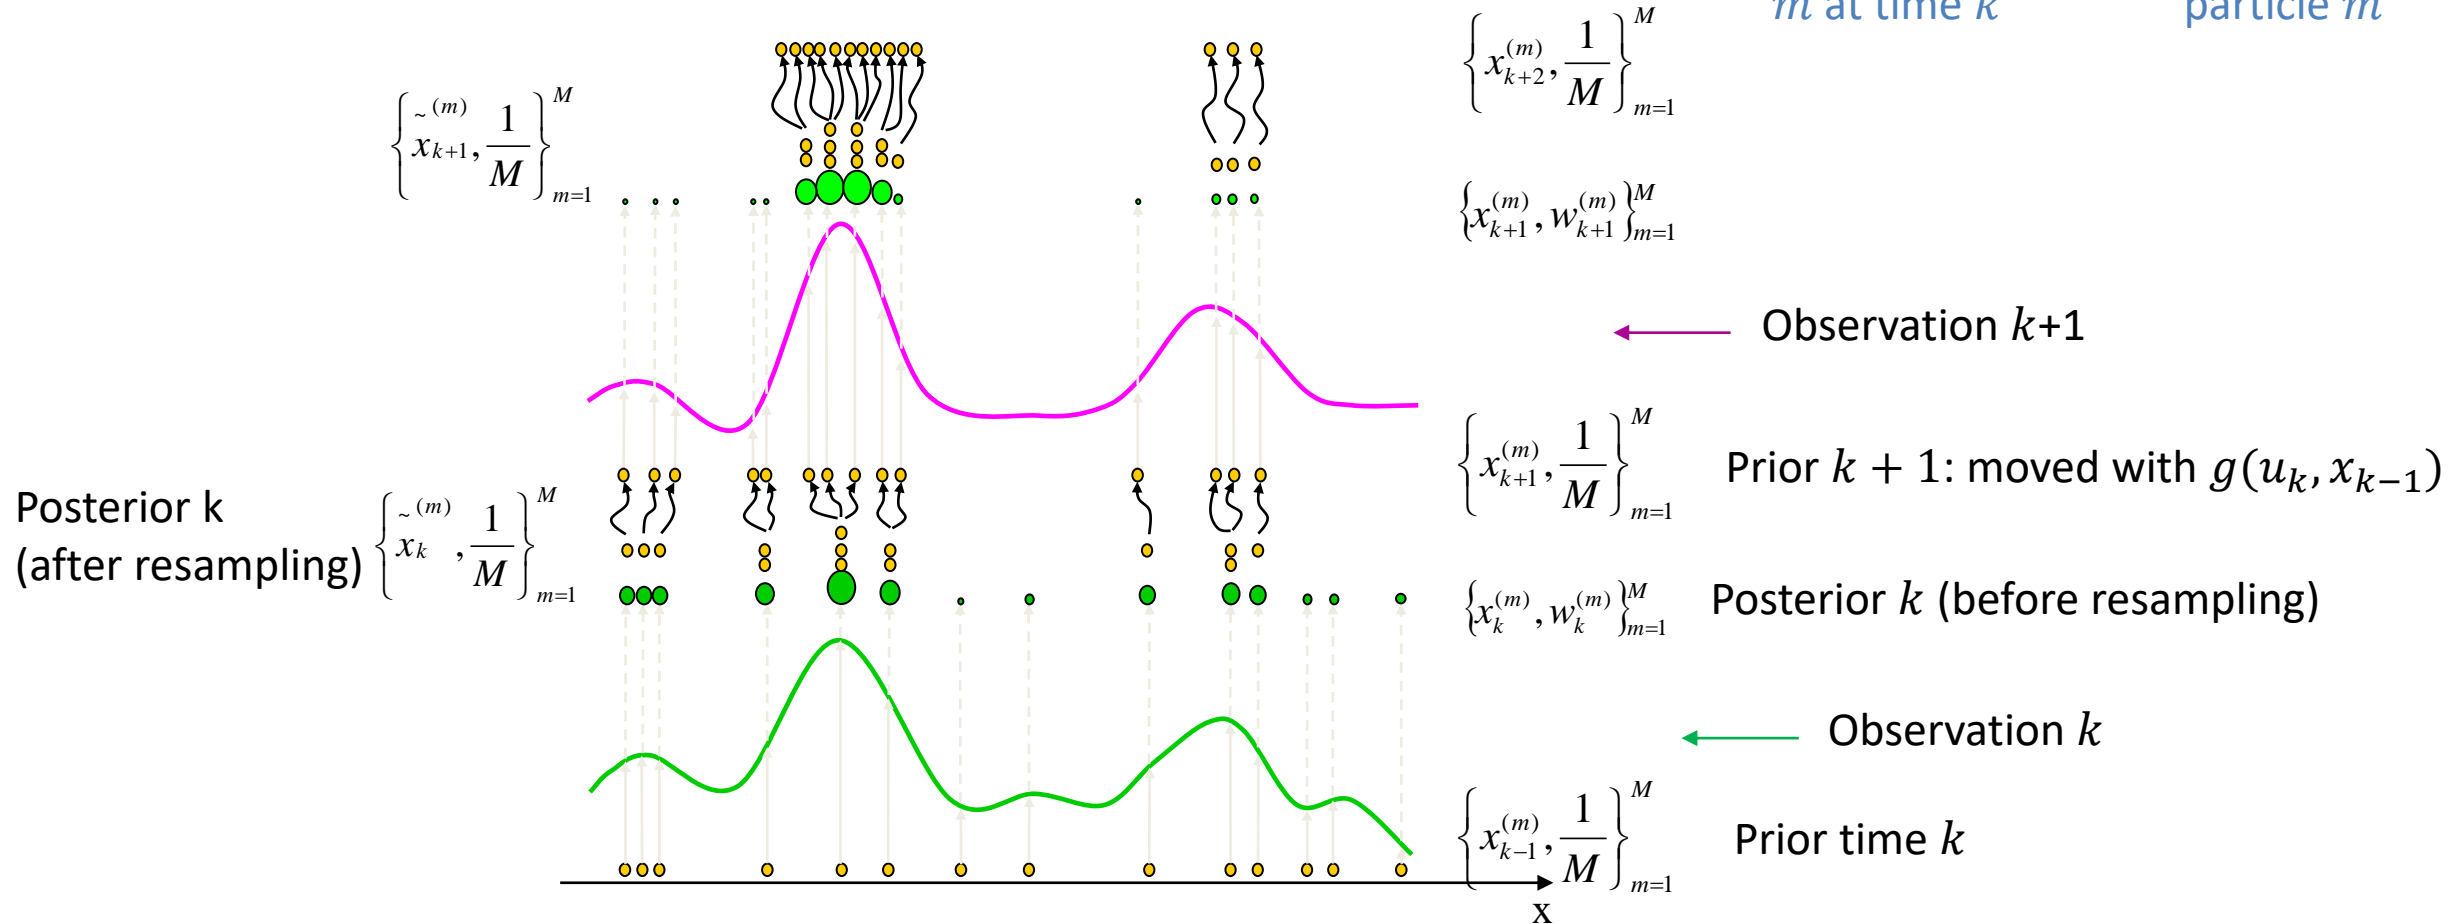


Illustration of the PF in 1D: SIR implementation, M particles: $\left\{x_{k-1}^{(m)}, w^{(m)}\right\}_{m=1}^M$

Weights here are represented by the area of the circle

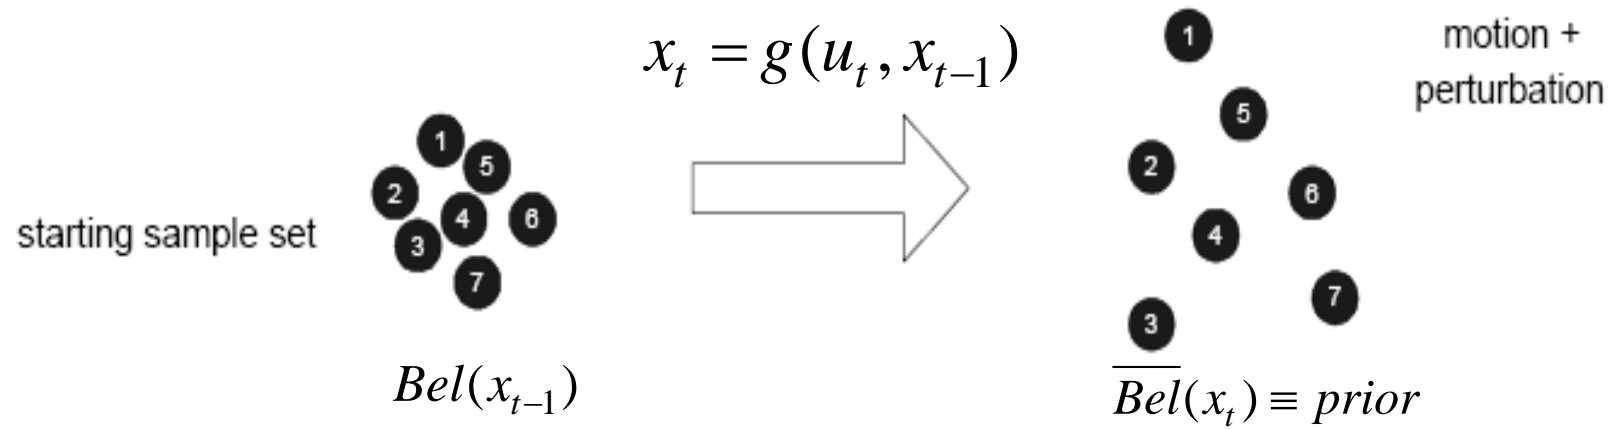
Pose of the particle m at time k Weight of the particle m



Particle Filter

1. Prediction


Implements the prediction $\overline{Bel}(g(u_t, x_{t-1}))$ in a sample-based form



```
% xP (3xn): Vector with the pose (x,y,θ) of the n particles
% U: covariance matrix 3x3 of the error motion
% Each particle p moves with the motion vector plus noise
for(p = 1:nParticles)
    xP(:,p) = compose(xP(:,p), u+sqrt(U)*randn(3,1));
end;
```

Particle Filter

2. Update. a) Observation prediction



```
% (1) Compute predicted observation for each particle x
zPred = DoObservationModel(xP(:,x), iFeature, Map);

% (2) Compute innovation
Innov = z - zPred;

% (3) Get likelihood (new weight)
L(x) = exp(-0.5 * Innov' * inv(REst) * Innov) + 0.001;
```

2. Update. b) Resampling based on weights (survival of the fittest)


IDEA: Particles with big weights will occupy a greater percentage of the y axis in a Likelihood accumulated Look-up table (LAT)

b1) Build a Look-up table for the Likelihood accumulation

Normalized Likelihood
accumulated table


$$\text{NLAT}(k) = \frac{1}{\text{sum}} \sum_{i=1}^k L(x_i)$$

$$\text{sum} = \sum_{i=1}^{np} L(x_i)$$



2. Update. b) Resampling based on weights (survival of the fittest)

b2) Randomly generate new particles from the weights



```
%LAT(k), table with accumulated likelihood
NLAT = LAT/sum;
%Randomly (uniform) choose a value y in [0,1] for all the particles
iSelect = rand(nParticles,1);
%Find particle that corresponds to each y value (just a look up)
iNextGeneration = interp1(NLAT,1:nParticles,iSelect,'nearest','extrap');
%copy selected particles for next generation..
xP = xP(:,iNextGeneration);
```


Example in 1D: Position tracking. A robot moves in a corridor with a range sensor that gets the distance to **ONE landmark** at the end of the corridor

PF is run every 5 m., starting at 0

Number of particles = 20

Sensor: distance to the landmark (+ gaussian noise)

— prior
— Posterior before resampling
— Observation likelihood




Example in 1D: Global localization

2 Landmarks, 1 observation

Number of particles = 500

- prior
- posterior
- Gaussian Observation likelihood to landmark 1
- Gaussian Observation likelihood to landmark 2

Given a measurement of $z = 12$ m we have to main hypothesis of the pose



Particle Filter

Advantages

- Model any probability distribution
- Model multimodal distributions (multiple hypothesis)
- Very simple implementation

Disadvantages

- Not as accurate as continuous random variable (it's an approximation). The more particles the more accurate
- Sometimes the PF fails: degeneracy problems
- Computational complexity: assumable for states of small dimension (3D), e.g. localization in 2D